

Adam Optimizer

February 3, 2024

1 Adam Optimizer from Scratch in Python

By Cristian Leo

Link Article: <https://medium.com/towards-data-science/the-math-behind-adam-optimizer-c41407efe59b>



1.1 Import Required Libraries

```
[1]: # Basic Libraries
import numpy as np
import pandas as pd

# Load Data
from sklearn.datasets import load_diabetes

# Data Visualization
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Model Fine Tuning
import optuna

# Filter Warnings
import warnings
warnings.filterwarnings('ignore')
```

/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages/tqdm/auto.py:21: TqdmWarning: IProgress not found. Please update jupyter and ipywidgets. See https://ipywidgets.readthedocs.io/en/stable/user_install.html
 from .autonotebook import tqdm as notebook_tqdm

1.2 Adam Optimizer Class

```
[2]: # Adam Optimizer (use the class from the previous response)
class AdamOptimizer:
    def __init__(self, learning_rate=0.001, beta1=0.9, beta2=0.999,
        ↪epsilon=1e-8):
        """
        Constructor for the AdamOptimizer class.

        Parameters
        -----
        learning_rate : float
            Learning rate for the optimizer.
        beta1 : float
            Exponential decay rate for the first moment estimates.
        beta2 : float
            Exponential decay rate for the second moment estimates.
        epsilon : float
            Small value to prevent division by zero.

        Returns
        -----
        None.
        """
        self.learning_rate = learning_rate
        self.beta1 = beta1
        self.beta2 = beta2
        self.epsilon = epsilon
        self.m = None
        self.v = None
        self.t = 0

    def initialize_moments(self, params):
```

```

    """
    Initializes the first and second moment estimates.

    Parameters
    -----
    params : dict
        Dictionary containing the model parameters.

    Returns
    -----
    None.
    """
    self.m = {k: np.zeros_like(v) for k, v in params.items()}
    self.v = {k: np.zeros_like(v) for k, v in params.items()}

def update_params(self, params, grads):
    """
    Updates the model parameters using the Adam optimizer.

    Parameters
    -----
    params : dict
        Dictionary containing the model parameters.
    grads : dict
        Dictionary containing the gradients for each parameter.

    Returns
    -----
    updated_params : dict
        Dictionary containing the updated model parameters.
    """
    if self.m is None or self.v is None:
        self.initialize_moments(params)

    self.t += 1
    updated_params = {}

    for key in params.keys():
        self.m[key] = self.beta1 * self.m[key] + (1 - self.beta1) * ↵
        ↵grads[key]
        self.v[key] = self.beta2 * self.v[key] + (1 - self.beta2) * np.
        ↵square(grads[key])

        m_corrected = self.m[key] / (1 - self.beta1 ** self.t)
        v_corrected = self.v[key] / (1 - self.beta2 ** self.t)

```

```

        updated_params[key] = params[key] - self.learning_rate *  $\frac{m\_corrected}{\sqrt{v\_corrected} + \epsilon}$ 
    return updated_params

```

1.3 Linear Regression Class

```

[3]: # Linear Regression Model
class LinearRegression:
    def __init__(self, n_features):
        """
        Constructor for the LinearRegression class.

        Parameters
        -----
        n_features : int
            Number of features in the input data.

        Returns
        -----
        None.
        """
        self.weights = np.random.randn(n_features)
        self.bias = np.random.randn()

    def predict(self, X):
        """
        Predicts the target variable given the input data.

        Parameters
        -----
        X : numpy array
            Input data.

        Returns
        -----
        numpy array
            Predictions.
        """
        return np.dot(X, self.weights) + self.bias

```

1.4 Model Trainer Class

```

[36]: class ModelTrainer:
    def __init__(self, model, optimizer, n_epochs):
        """

```

```

    Constructor for the ModelTrainer class.

    Parameters
    -----
    model : object
        Model to be trained.
    optimizer : object
        Optimizer to be used for training.
    n_epochs : int
        Number of training epochs.

    Returns
    -----
    None.
    """
    self.model = model
    self.optimizer = optimizer
    self.n_epochs = n_epochs

def compute_gradients(self, X, y):
    """
    Computes the gradients of the mean squared error loss function
    with respect to the model parameters.

    Parameters
    -----
    X : numpy array
        Input data.
    y : numpy array
        Target variable.

    Returns
    -----
    dict
        Dictionary containing the gradients for each parameter.
    """
    predictions = self.model.predict(X)
    errors = predictions - y
    dW = 2 * np.dot(X.T, errors) / len(y)
    db = 2 * np.mean(errors)
    return {'weights': dW, 'bias': db}

def train(self, X, y, verbose=False):
    """
    Runs the training loop, updating the model parameters and optionally
    printing the loss.

```

```

Parameters
-----
X : numpy array
    Input data.
y : numpy array
    Target variable.

Returns
-----
None.
"""
for epoch in range(self.n_epochs):
    grads = self.compute_gradients(X, y)
    params = {'weights': self.model.weights, 'bias': self.model.bias}
    updated_params = self.optimizer.update_params(params, grads)

    self.model.weights = updated_params['weights']
    self.model.bias = updated_params['bias']

    # Optionally, print loss here to observe training
    loss = np.mean((self.model.predict(X) - y) ** 2)
    if epoch % 1000 == 0 and verbose:
        print(f"Epoch {epoch}, Loss: {loss}")

```

1.5 Load Diabetes Data

```

[16]: # Load the diabetes dataset
diabetes = load_diabetes()
df = pd.DataFrame(data=diabetes.data, columns=diabetes.feature_names)
df['target'] = diabetes.target
df.head()

```

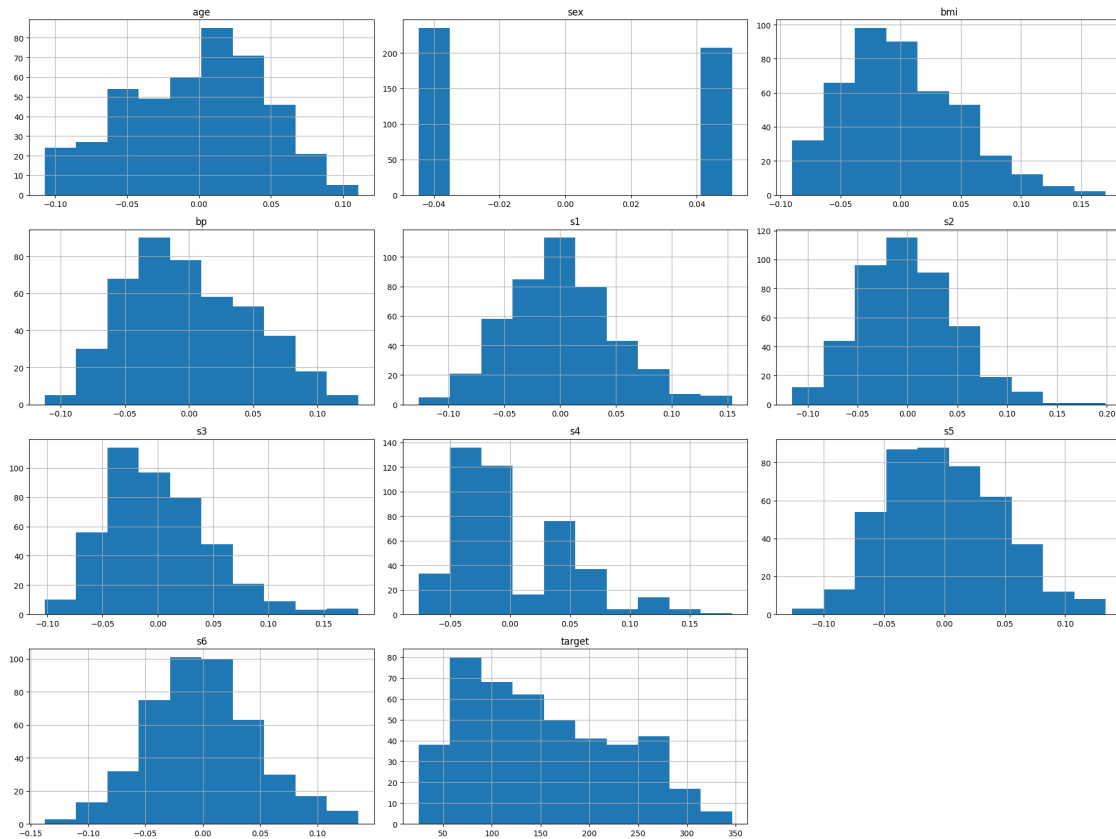
```

[16]:
   age      sex      bmi      bp      s1      s2      s3 \
0  0.038076  0.050680  0.061696  0.021872 -0.044223 -0.034821 -0.043401
1 -0.001882 -0.044642 -0.051474 -0.026328 -0.008449 -0.019163  0.074412
2  0.085299  0.050680  0.044451 -0.005670 -0.045599 -0.034194 -0.032356
3 -0.089063 -0.044642 -0.011595 -0.036656  0.012191  0.024991 -0.036038
4  0.005383 -0.044642 -0.036385  0.021872  0.003935  0.015596  0.008142

      s4      s5      s6  target
0 -0.002592  0.019907 -0.017646   151.0
1 -0.039493 -0.068332 -0.092204    75.0
2 -0.002592  0.002861 -0.025930   141.0
3  0.034309  0.022688 -0.009362   206.0
4 -0.002592 -0.031988 -0.046641   135.0

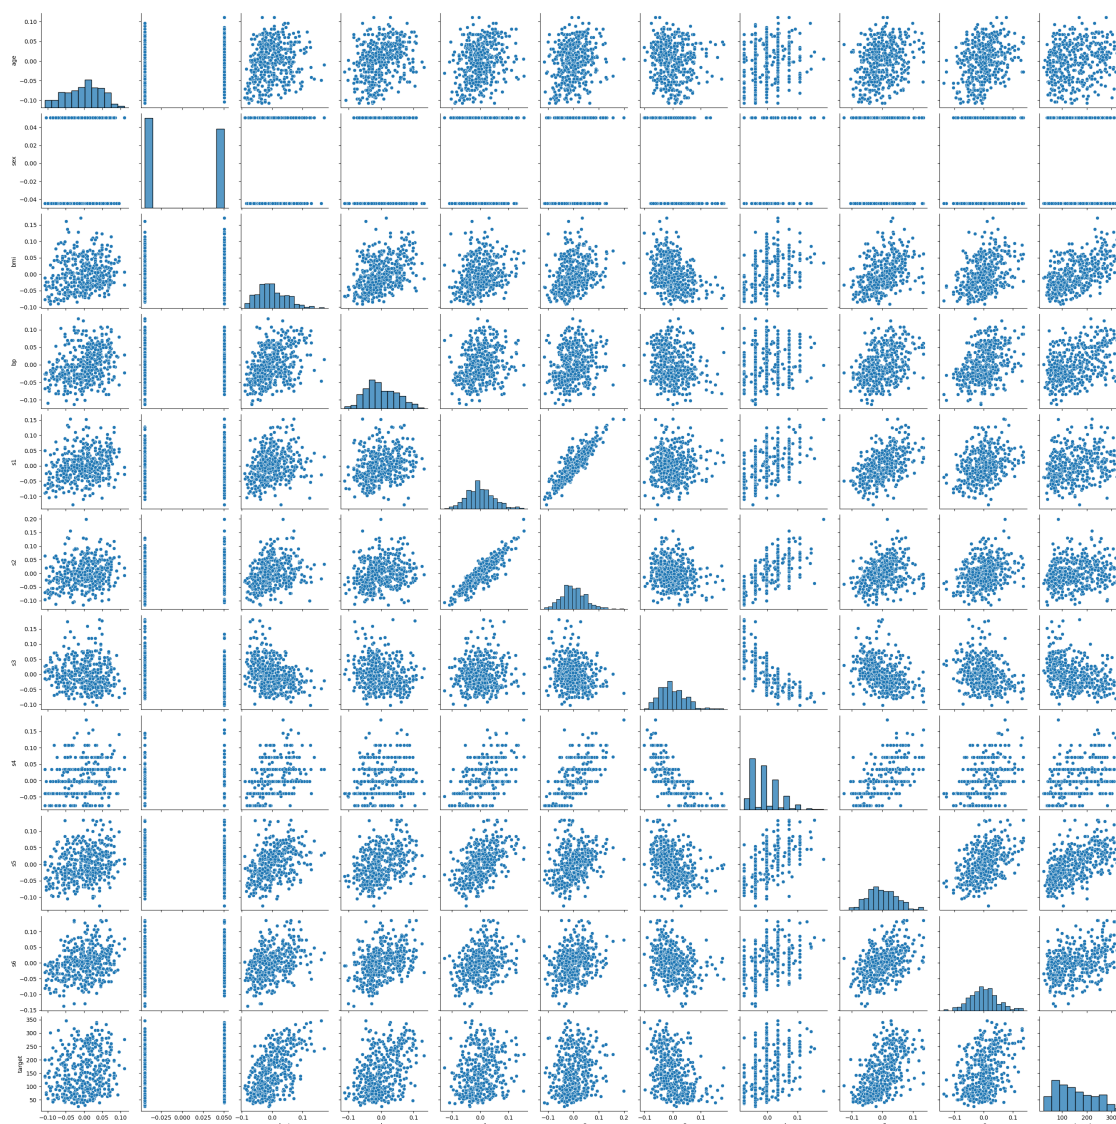
```

```
[17]: # Create histograms for each feature
df.hist(bins=10, figsize=(20, 15))
plt.tight_layout()
plt.show()
```



```
[18]: # Pairplot of the features
sns.pairplot(df)
```

```
[18]: <seaborn.axisgrid.PairGrid at 0x29ed37500>
```



1.6 Split Data

```
[19]: # Get the input features (X) and target values (y)
X = diabetes.data
y = diabetes.target

# Split the dataset into training and test sets
def split_dataset(X, y, test_ratio=0.2):
    indices = np.random.permutation(len(X))
    test_size = int(len(X) * test_ratio)
    test_indices = indices[:test_size]
    train_indices = indices[test_size:]
    return X[train_indices], X[test_indices], y[train_indices], y[test_indices]
```



```
X_train, X_test, y_train, y_test = split_dataset(X, y)
X_train, X_val, y_train, y_val = split_dataset(X_train, y_train)
```

1.7 Fine Tune Model with Optuna

```
[20]: y_train.shape
```

```
[20]: (284,)
```

```
[37]: def objective(trial):
    n_features = X_train.shape[1]

    learning_rate = trial.suggest_loguniform('learning_rate', 1e-5, 1e-1)
    beta1 = trial.suggest_uniform('beta1', 0.9, 0.999)
    beta2 = trial.suggest_uniform('beta2', 0.99, 0.9999)
    epsilon = trial.suggest_loguniform('epsilon', 1e-10, 1e-5)

    n_epochs = trial.suggest_int('epochs', 1000, 100000)

    # Define the model
    model = LinearRegression(n_features)
    optimizer = AdamOptimizer(learning_rate=learning_rate, beta1=beta1,
    ↪beta2=beta2, epsilon=epsilon)
    trainer = ModelTrainer(model, optimizer, n_epochs=n_epochs)

    # Train the model
    trainer.train(X_train, y_train, verbose=False)

    # Compute the validation loss
    val_loss = np.mean((model.predict(X_val) - y_val) ** 2)

    return val_loss

    # Create the study object
    optuna.logging.set_verbosity(optuna.logging.WARNING)
    study = optuna.create_study(direction='minimize', sampler=optuna.samplers.
    ↪TPESampler(seed=42))

    # Optimize the study, use more trials to obtain better results, use less trials
    ↪to be more cost-efficient
    study.optimize(objective, n_trials=10)
```

```
[38]: # Print optimization results
print("-----")
print('Number of finished trials:', len(study.trials))
```

```

print('Best trial:')
for key, value in study.best_trial.params.items():
    if key == 'epochs' or key == 'epsilon':
        print(f'    {key}: {value}')
    else:
        print(f'    {key}: {value:.3f}')

```

```

-----
Number of finished trials: 10
Best trial:
    learning_rate: 0.017
    beta1: 0.930
    beta2: 0.991
    epsilon: 2.637333993381524e-07
    epochs: 44575

```

```

[40]: # Get the best model
n_features = X_train.shape[1]
best_model = LinearRegression(n_features)
optimizer = AdamOptimizer(learning_rate=study.best_params['learning_rate'],
                           beta1=study.best_params['beta1'],
                           beta2=study.best_params['beta2'],
                           epsilon=study.best_params['epsilon'])

# Train the model
trainer = ModelTrainer(best_model, optimizer, n_epochs=study.
    ↪best_params['epochs'])
trainer.train(X_train, y_train)

# Compute the test loss
test_loss = np.mean((best_model.predict(X_test) - y_test) ** 2)**0.5
print(f'Test loss: {test_loss:.2f}')

```

Test loss: 53.55

```

[31]: # Plot the predictions vs the actual values
plt.figure(figsize=(10, 10))
plt.scatter(y_test, best_model.predict(X_test), c='crimson', alpha=0.7)
plt.title('Predictions vs Actual Values', fontsize=16)
plt.xlabel('Actual target values')
plt.ylabel('Predicted target values')
plt.show()

```

