

# lstm

April 29, 2024

## 1 LSTM From Scratch in Python

*By Cristian Leo*



```
[1]: import numpy as np
import requests
import pandas as pd
import matplotlib.pyplot as plt

# Custom classes (built from scratch)
from src.model import WeightInitializer
from src.trainer import PlotManager, EarlyStopping
```

```
[2]: class LSTM:
    """
    Long Short-Term Memory (LSTM) network.

    Parameters:
```

```

- input_size: int, dimensionality of input space
- hidden_size: int, number of LSTM units
- output_size: int, dimensionality of output space
- init_method: str, weight initialization method (default: 'xavier')
"""
def __init__(self, input_size, hidden_size, output_size,
init_method='xavier'):
    self.input_size = input_size
    self.hidden_size = hidden_size
    self.output_size = output_size
    self.weight_initializer = WeightInitializer(method=init_method)

    # Initialize weights
    self.wf = self.weight_initializer.initialize((hidden_size, hidden_size,
↪ input_size))
    self.wi = self.weight_initializer.initialize((hidden_size, hidden_size,
↪ input_size))
    self.wo = self.weight_initializer.initialize((hidden_size, hidden_size,
↪ input_size))
    self.wc = self.weight_initializer.initialize((hidden_size, hidden_size,
↪ input_size))

    # Initialize biases
    self.bf = np.zeros((hidden_size, 1))
    self.bi = np.zeros((hidden_size, 1))
    self.bo = np.zeros((hidden_size, 1))
    self.bc = np.zeros((hidden_size, 1))

    # Initialize output layer weights and biases
    self.why = self.weight_initializer.initialize((output_size,
↪ hidden_size))
    self.by = np.zeros((output_size, 1))

    @staticmethod
    def sigmoid(z):
        """
        Sigmoid activation function.

        Parameters:
        - z: np.ndarray, input to the activation function

        Returns:
        - np.ndarray, output of the activation function
        """
        return 1 / (1 + np.exp(-z))

```

```

@staticmethod
def dsigmoid(y):
    """
    Derivative of the sigmoid activation function.

    Parameters:
    - y: np.ndarray, output of the sigmoid activation function

    Returns:
    - np.ndarray, derivative of the sigmoid function
    """
    return y * (1 - y)

@staticmethod
def dtanh(y):
    """
    Derivative of the hyperbolic tangent activation function.

    Parameters:
    - y: np.ndarray, output of the hyperbolic tangent activation function

    Returns:
    - np.ndarray, derivative of the hyperbolic tangent function
    """
    return 1 - y * y

def forward(self, x):
    """
    Forward pass through the LSTM network.

    Parameters:
    - x: np.ndarray, input to the network

    Returns:
    - np.ndarray, output of the network
    - list, caches containing intermediate values for backpropagation
    """
    caches = []
    h_prev = np.zeros((self.hidden_size, 1))
    c_prev = np.zeros((self.hidden_size, 1))
    h = h_prev
    c = c_prev

    for t in range(x.shape[0]):
        x_t = x[t].reshape(-1, 1)
        combined = np.vstack((h_prev, x_t))

```

```

        f = self.sigmoid(np.dot(self.wf, combined) + self.bf)
        i = self.sigmoid(np.dot(self.wi, combined) + self.bi)
        o = self.sigmoid(np.dot(self.wo, combined) + self.bo)
        c_ = np.tanh(np.dot(self.wc, combined) + self.bc)

        c = f * c_prev + i * c_
        h = o * np.tanh(c)

        cache = (h_prev, c_prev, f, i, o, c_, x_t, combined, c, h)
        caches.append(cache)

    h_prev, c_prev = h, c

    y = np.dot(self.why, h) + self.by
    return y, caches

def backward(self, dy, caches, clip_value=1.0):
    """
    Backward pass through the LSTM network.

    Parameters:
    - dy: np.ndarray, gradient of the loss with respect to the output
    - caches: list, caches from the forward pass
    - clip_value: float, value to clip gradients to (default: 1.0)

    Returns:
    - tuple, gradients of the loss with respect to the parameters
    """
    dWf, dWi, dWo, dWc = [np.zeros_like(w) for w in (self.wf, self.wi, self.
    ↪wo, self.wc)]
    dbf, dbi, dbo, dbc = [np.zeros_like(b) for b in (self.bf, self.bi, self.
    ↪bo, self.bc)]
    dWhy = np.zeros_like(self.why)
    dby = np.zeros_like(self.by)

    # Ensure dy is reshaped to match output size
    dy = dy.reshape(self.output_size, -1)
    dh_next = np.zeros((self.hidden_size, 1)) # shape must match ↪
    ↪hidden_size
    dc_next = np.zeros_like(dh_next)

    for cache in reversed(caches):
        h_prev, c_prev, f, i, o, c_, x_t, combined, c, h = cache

        # Add gradient from next step to current output gradient
        dh = np.dot(self.why.T, dy) + dh_next
        dc = dc_next + (dh * o * self.dtanh(np.tanh(c)))

```

```

df = dc * c_prev * self.dsigmoid(f)
di = dc * c_ * self.dsigmoid(i)
do = dh * self.dtanh(np.tanh(c))
dc_ = dc * i * self.dtanh(c_)

dcombined_f = np.dot(self.wf.T, df)
dcombined_i = np.dot(self.wi.T, di)
dcombined_o = np.dot(self.wo.T, do)
dcombined_c = np.dot(self.wc.T, dc_)

dcombined = dcombined_f + dcombined_i + dcombined_o + dcombined_c
dh_next = dcombined[:self.hidden_size]
dc_next = f * dc

dWf += np.dot(df, combined.T)
dWi += np.dot(di, combined.T)
dWo += np.dot(do, combined.T)
dWc += np.dot(dc_, combined.T)

dbf += df.sum(axis=1, keepdims=True)
dbi += di.sum(axis=1, keepdims=True)
dbo += do.sum(axis=1, keepdims=True)
dbc += dc_.sum(axis=1, keepdims=True)

dWhy += np.dot(dy, h.T)
dby += dy

gradients = (dWf, dWi, dWo, dWc, dbf, dbi, dbo, dbc, dWhy, dby)

# Gradient clipping
for i in range(len(gradients)):
    np.clip(gradients[i], -clip_value, clip_value, out=gradients[i])

return gradients

def update_params(self, grads, learning_rate):
    """
    Update the parameters of the network using the gradients.

    Parameters:
    - grads: tuple, gradients of the loss with respect to the parameters
    - learning_rate: float, learning rate
    """
    dWf, dWi, dWo, dWc, dbf, dbi, dbo, dbc, dWhy, dby = grads

    self.wf -= learning_rate * dWf

```

```

self.wi -= learning_rate * dWi
self.wo -= learning_rate * dWo
self.wc -= learning_rate * dWc

self.bf -= learning_rate * dbf
self.bi -= learning_rate * dbi
self.bo -= learning_rate * dbo
self.bc -= learning_rate * dbc

self.why -= learning_rate * dWhy
self.by -= learning_rate * dby

```

```

[3]: class LSTMTrainer:
    """
    Trainer for the LSTM network.

    Parameters:
    - model: LSTM, the LSTM network to train
    - learning_rate: float, learning rate for the optimizer
    - patience: int, number of epochs to wait before early stopping
    - verbose: bool, whether to print training information
    - delta: float, minimum change in validation loss to qualify as an
    ↪improvement
    """
    def __init__(self, model, learning_rate=0.01, patience=7, verbose=True,
    ↪delta=0):
        self.model = model
        self.learning_rate = learning_rate
        self.train_losses = []
        self.val_losses = []
        self.early_stopping = EarlyStopping(patience, verbose, delta)

    def train(self, X_train, y_train, X_val=None, y_val=None, epochs=10,
    ↪batch_size=1, clip_value=1.0):
        """
        Train the LSTM network.

        Parameters:
        - X_train: np.ndarray, training data
        - y_train: np.ndarray, training labels
        - X_val: np.ndarray, validation data
        - y_val: np.ndarray, validation labels
        - epochs: int, number of training epochs
        - batch_size: int, size of mini-batches
        - clip_value: float, value to clip gradients to
        """
        for epoch in range(epochs):

```

```

epoch_losses = []
for i in range(0, len(X_train), batch_size):
    batch_X = X_train[i:i + batch_size]
    batch_y = y_train[i:i + batch_size]
    losses = []

    for x, y_true in zip(batch_X, batch_y):
        y_pred, caches = self.model.forward(x)
        loss = self.compute_loss(y_pred, y_true.reshape(-1, 1))
        losses.append(loss)

        # Backpropagation to get gradients
        dy = y_pred - y_true.reshape(-1, 1)
        grads = self.model.backward(dy, caches, clip_value=clip_value)
        self.model.update_params(grads, self.learning_rate)

    batch_loss = np.mean(losses)
    epoch_losses.append(batch_loss)

avg_epoch_loss = np.mean(epoch_losses)
self.train_losses.append(avg_epoch_loss)

if X_val is not None and y_val is not None:
    val_loss = self.validate(X_val, y_val)
    self.val_losses.append(val_loss)

    if epoch % 10 == 0:
        print(f'Epoch {epoch + 1}/{epochs} - Loss: {avg_epoch_loss:.5f}, Val Loss: {val_loss:.5f}')

        # Check early stopping condition
        self.early_stopping(val_loss)
        if self.early_stopping.early_stop:
            print("Early stopping")
            break
    else:
        print(f'Epoch {epoch + 1}/{epochs} - Loss: {avg_epoch_loss:.5f}')

def compute_loss(self, y_pred, y_true):
    """
    Compute mean squared error loss.
    """
    return np.mean((y_pred - y_true) ** 2)

```

```

def validate(self, X_val, y_val):
    """
    Validate the model on a separate set of data.
    """
    val_losses = []
    for x, y_true in zip(X_val, y_val):
        y_pred, _ = self.model.forward(x)
        loss = self.compute_loss(y_pred, y_true.reshape(-1, 1))
        val_losses.append(loss)
    return np.mean(val_losses)

```

```

[4]: class TimeSeriesDataset:
    """
    Dataset class for time series data.

    Parameters:
    - ticker: str, stock ticker symbol
    - start_date: str, start date for data retrieval
    - end_date: str, end date for data retrieval
    - look_back: int, number of previous time steps to include in each sample
    - train_size: float, proportion of data to use for training
    """

    def __init__(self, ticker, start_date, end_date, look_back=1, train_size=0.
↪67, api_key=None):
        self.ticker = ticker
        self.start_date = start_date
        self.end_date = end_date
        self.look_back = look_back
        self.train_size = train_size
        self.api_key = api_key

    def load_data(self):
        """
        Load stock data using AlphaVantage API.

        Returns:
        - np.ndarray, training data
        - np.ndarray, testing data
        """
        if self.api_key is None:
            df = pd.read_csv('data/google.csv')
        else:
            url = f'https://www.alphavantage.co/query?
↪function=TIME_SERIES_DAILY&symbol={self.ticker}&outputsize=full&apikey={self.
↪api_key}'
            response = requests.get(url)
            data = response.json()

```



```

        df = pd.DataFrame(data['Time Series (Daily)']).T
        df.to_csv('data/google.csv')
    df = df.sort_index()
    df = df.loc[self.start_date:self.end_date]
    df = df[['4. close']].astype(float) # Use closing price
    df = self.MinMaxScaler(df.values) # Convert DataFrame to numpy array
    train_size = int(len(df) * self.train_size)
    train, test = df[0:train_size,:], df[train_size:len(df),:]
    return train, test

def MinMaxScaler(self, data):
    """
    Min-max scaling of the data.

    Parameters:
    - data: np.ndarray, input data
    """
    numerator = data - np.min(data, 0)
    denominator = np.max(data, 0) - np.min(data, 0)
    return numerator / (denominator + 1e-7)

def create_dataset(self, dataset):
    """
    Create the dataset for time series prediction.

    Parameters:
    - dataset: np.ndarray, input data

    Returns:
    - np.ndarray, input data
    - np.ndarray, output data
    """
    dataX, dataY = [], []
    for i in range(len(dataset)-self.look_back):
        a = dataset[i:(i + self.look_back), 0]
        dataX.append(a)
        dataY.append(dataset[i + self.look_back, 0])
    return np.array(dataX), np.array(dataY)

def get_train_test(self):
    """
    Get the training and testing data.

    Returns:
    - np.ndarray, training input
    - np.ndarray, training output
    - np.ndarray, testing input

```

```

- np.ndarray, testing output
"""
train, test = self.load_data()
trainX, trainY = self.create_dataset(train)
testX, testY = self.create_dataset(test)
return trainX, trainY, testX, testY

```

```

[5]: # Instantiate the dataset
dataset = TimeSeriesDataset('GOOGL', '2010-1-1', '2023-12-31', look_back=1,
    ↪ train_size=0.7, api_key=None)
trainX, trainY, testX, testY = dataset.get_train_test()

# Plot the data
# Combine train and test data
combined = np.concatenate((trainY, testY))

# Plot the data
plt.figure(figsize=(14, 5))
plt.plot(combined, label='Google Stock Price', linewidth=2, color='dodgerblue')
plt.title('Google Stock Price', fontsize=20)
plt.xlabel('Time', fontsize=16)
plt.ylabel('Normalized Stock Price', fontsize=16)
plt.grid(True)
plt.legend(fontsize=14)
plt.xticks(fontsize=12)
plt.yticks(fontsize=12)
plt.show()

```



```

[6]: # Reshape input to be [samples, time steps, features]
trainX = np.reshape(trainX, (trainX.shape[0], trainX.shape[1], 1))
testX = np.reshape(testX, (testX.shape[0], testX.shape[1], 1))

```

```

look_back = 1 # Number of previous time steps to include in each sample
hidden_size = 256 # Number of LSTM units
output_size = 1 # Dimensionality of the output space

lstm = LSTM(input_size=1, hidden_size=hidden_size, output_size=output_size)

# Create and train the LSTM using LSTMTrainer
trainer = LSTMTrainer(lstm, learning_rate=1e-3, patience=50, verbose=True,
    ↪delta=0.001)
trainer.train(trainX, trainY, testX, testY, epochs=1000, batch_size=32)

```

```

Epoch 1/1000 - Loss: 0.24601, Val Loss: 0.41803
Epoch 11/1000 - Loss: 0.06322, Val Loss: 0.05970
Epoch 21/1000 - Loss: 0.05178, Val Loss: 0.02055
Epoch 31/1000 - Loss: 0.04755, Val Loss: 0.01107
Epoch 41/1000 - Loss: 0.04429, Val Loss: 0.00667
Epoch 51/1000 - Loss: 0.04169, Val Loss: 0.00395
Epoch 61/1000 - Loss: 0.03962, Val Loss: 0.00216
Epoch 71/1000 - Loss: 0.03797, Val Loss: 0.00103
Epoch 81/1000 - Loss: 0.03666, Val Loss: 0.00037
Epoch 91/1000 - Loss: 0.03563, Val Loss: 0.00006
Epoch 101/1000 - Loss: 0.03480, Val Loss: 0.00001
Epoch 111/1000 - Loss: 0.03415, Val Loss: 0.00014
Epoch 121/1000 - Loss: 0.03362, Val Loss: 0.00040
Early stopping

```

```

[7]: plot_manager = PlotManager()

# Inside your training loop
plot_manager.plot_losses(trainer.train_losses, trainer.val_losses)

# After your training loop
plot_manager.show_plots()

```

