


NN

April 2, 2024

1 Neural Network From Scratch in Python

By Cristian Leo



1.1 Libraries

```
[1]: import numpy as np # linear algebra

import matplotlib.pyplot as plt # this is used for the plot the graph

from sklearn.datasets import load_digits # this is used for import the dataset
from sklearn.model_selection import train_test_split # to split the data into
    ↳ two parts
from sklearn.preprocessing import MinMaxScaler # for normalization
from sklearn.preprocessing import OneHotEncoder # for one hot encoding

import optuna # for hyperparameter tuning

import warnings # this is used for ignoring the warnings
```

```
warnings.filterwarnings("ignore")
```

```
/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-  
packages/tqdm/auto.py:21: TqdmWarning: IProgress not found. Please update  
jupyter and ipywidgets. See  
https://ipywidgets.readthedocs.io/en/stable/user\_install.html  
from .autonotebook import tqdm as notebook_tqdm
```

1.2 Neural Network Class

```
[2]: class NeuralNetwork:  
    """  
    A simple neural network with one hidden layer.  
  
    Parameters:  
    -----  
    input_size: int  
        The number of input features  
    hidden_size: int  
        The number of neurons in the hidden layer  
    output_size: int  
        The number of neurons in the output layer  
    loss_func: str  
        The loss function to use. Options are 'mse' for mean squared error,   
    ↪ 'log_loss' for logistic loss, and 'categorical_crossentropy' for categorical   
    ↪ crossentropy.  
    """  
    def __init__(self, input_size, hidden_size, output_size, loss_func='mse'):  
        self.input_size = input_size  
        self.hidden_size = hidden_size  
        self.output_size = output_size  
        self.loss_func = loss_func  
  
        # Initialize weights and biases  
        self.weights1 = np.random.randn(self.input_size, self.hidden_size)  
        self.bias1 = np.zeros((1, self.hidden_size))  
        self.weights2 = np.random.randn(self.hidden_size, self.output_size)  
        self.bias2 = np.zeros((1, self.output_size))  
  
        # track loss  
        self.train_loss = []  
        self.test_loss = []  
  
    def __str__(self):  
        """  
        Print the neural network architecture.  
        """
```

```

        return f"Neural Network Layout:\nInput Layer: {self.input_size}\n
↳neurons\nHidden Layer: {self.hidden_size} neurons\nOutput Layer: {self.
↳output_size} neurons\nLoss Function: {self.loss_func}"

def forward(self, X):
    """
    Perform forward propagation.

    Parameters:
    -----
    X: numpy array
        The input data

    Returns:
    -----
    numpy array
        The predicted output
    """
    # Perform forward propagation
    self.z1 = np.dot(X, self.weights1) + self.bias1
    self.a1 = self.sigmoid(self.z1)
    self.z2 = np.dot(self.a1, self.weights2) + self.bias2
    if self.loss_func == 'categorical_crossentropy':
        self.a2 = self.softmax(self.z2)
    else:
        self.a2 = self.sigmoid(self.z2)
    return self.a2

def backward(self, X, y, learning_rate):
    """
    Perform backpropagation.

    Parameters:
    -----
    X: numpy array
        The input data
    y: numpy array
        The target output
    learning_rate: float
        The learning rate
    """
    # Perform backpropagation
    m = X.shape[0]

    # Calculate gradients
    if self.loss_func == 'mse':
        self.dz2 = self.a2 - y

```

```

elif self.loss_func == 'log_loss':
    self.dz2 = -(y/self.a2 - (1-y)/(1-self.a2))
elif self.loss_func == 'categorical_crossentropy':
    self.dz2 = self.a2 - y
else:
    raise ValueError('Invalid loss function')

self.dw2 = (1 / m) * np.dot(self.a1.T, self.dz2)
self.db2 = (1 / m) * np.sum(self.dz2, axis=0, keepdims=True)
self.dz1 = np.dot(self.dz2, self.weights2.T) * self.
sigmoid_derivative(self.a1)
self.dw1 = (1 / m) * np.dot(X.T, self.dz1)
self.db1 = (1 / m) * np.sum(self.dz1, axis=0, keepdims=True)

# Update weights and biases
self.weights2 -= learning_rate * self.dw2
self.bias2 -= learning_rate * self.db2
self.weights1 -= learning_rate * self.dw1
self.bias1 -= learning_rate * self.db1

def sigmoid(self, x):
    """
    Sigmoid activation function.

    Parameters:
    -----
    x: numpy array
        The input data

    Returns:
    -----
    numpy array
        The output of the sigmoid function
    """
    return 1 / (1 + np.exp(-x))

def sigmoid_derivative(self, x):
    """
    Derivative of the sigmoid activation function.

    Parameters:
    -----
    x: numpy array
        The input data

    Returns:
    -----

```

```

    numpy array
    The output of the derivative of the sigmoid function
    """
    return x * (1 - x)

def softmax(self, x):
    """
    Softmax activation function.

    Parameters:
    -----
    x: numpy array
        The input data

    Returns:
    -----
    numpy array
        The output of the softmax function
    """
    exps = np.exp(x - np.max(x, axis=1, keepdims=True))
    return exps/np.sum(exps, axis=1, keepdims=True)

```

1.3 Trainer Class

```

[3]: class Trainer:
    """
    A class to train a neural network.

    Parameters:
    -----
    model: NeuralNetwork
        The neural network model to train
    loss_func: str
        The loss function to use. Options are 'mse' for mean squared error,
        ↪ 'log_loss' for logistic loss, and 'categorical_crossentropy' for categorical
        ↪ crossentropy.
    """
    def __init__(self, model, loss_func='mse'):
        self.model = model
        self.loss_func = loss_func
        self.train_loss = []
        self.test_loss = []

    def calculate_loss(self, y_true, y_pred):
        """
        Calculate the loss.

```

```

Parameters:
-----
y_true: numpy array
    The true output
y_pred: numpy array
    The predicted output

Returns:
-----
float
    The loss
"""
if self.loss_func == 'mse':
    return np.mean((y_pred - y_true)**2)
elif self.loss_func == 'log_loss':
    return -np.mean(y_true*np.log(y_pred) + (1-y_true)*np.log(1-y_pred))
elif self.loss_func == 'categorical_crossentropy':
    return -np.mean(y_true*np.log(y_pred))
else:
    raise ValueError('Invalid loss function')

def train(self, X_train, y_train, X_test, y_test, epochs, learning_rate):
    """
    Train the neural network.

    Parameters:
    -----
    X_train: numpy array
        The training input data
    y_train: numpy array
        The training target output
    X_test: numpy array
        The test input data
    y_test: numpy array
        The test target output
    epochs: int
        The number of epochs to train the model
    learning_rate: float
        The learning rate
    """
    for _ in range(epochs):
        self.model.forward(X_train)
        self.model.backward(X_train, y_train, learning_rate)
        train_loss = self.calculate_loss(y_train, self.model.a2)
        self.train_loss.append(train_loss)

    self.model.forward(X_test)

```

```
test_loss = self.calculate_loss(y_test, self.model.a2)
self.test_loss.append(test_loss)
```


1.4 Load Dataset

```
[4]: # Load the digits dataset
digits = load_digits()

# Plot the first 10 images
fig, axes = plt.subplots(2, 5, figsize=(10, 5))
axes = axes.ravel()

for i in range(10):
    axes[i].imshow(digits.images[i], cmap='gray')
    axes[i].axis('off')
    axes[i].set_title(f"Label: {digits.target[i]}")

plt.tight_layout()
plt.show()
```



1.5 Data Preprocessing

```
[5]: # Preprocess the dataset
scaler = MinMaxScaler()
X = scaler.fit_transform(digits.data)
y = digits.target

# One-hot encode the target output
encoder = OneHotEncoder(sparse=False)
```

```

y_onehot = encoder.fit_transform(y.reshape(-1, 1))

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y_onehot, test_size=0.2,
    random_state=42)

```

1.6 Create Neural Network

```

[6]: # Create an instance of the NeuralNetwork class
input_size = X.shape[1]
hidden_size = 64
output_size = len(np.unique(y))
loss_func = 'categorical_crossentropy'
epochs = 1000
learning_rate = 0.1

nn = NeuralNetwork(input_size, hidden_size, output_size, loss_func)

# Print the neural network architecture
print(nn)

```

Neural Network Layout:

Input Layer: 64 neurons

Hidden Layer: 64 neurons

Output Layer: 10 neurons

Loss Function: categorical_crossentropy

1.7 Train NN

```

[7]: trainer = Trainer(nn, loss_func)
trainer.train(X_train, y_train, X_test, y_test, epochs, learning_rate)

# Convert y_test from one-hot encoding to labels
y_test_labels = np.argmax(y_test, axis=1)

# Evaluate the performance of the neural network
predictions = np.argmax(nn.forward(X_test), axis=1)
accuracy = np.mean(predictions == y_test_labels)
print(f"Accuracy: {accuracy:.2%}")

```

Accuracy: 93.33%

1.8 Plot Loss


```

[8]: plt.plot(trainer.train_loss, label='Train Loss')
plt.plot(trainer.test_loss, label='Test Loss')
plt.title('Train and Test Loss')

```



```
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



1.9 Fine-Tune NN

```
[9]: def objective(trial):
    # Define hyperparameters
    hidden_size = trial.suggest_int('hidden_size', 32, 128)
    learning_rate = trial.suggest_loguniform('learning_rate', 1e-4, 1e-1)
    epochs = trial.suggest_int('epochs', 500, 10000)

    # Create and train the neural network
    nn = NeuralNetwork(input_size, hidden_size, output_size, loss_func)
    trainer = Trainer(nn, loss_func)
    trainer.train(X_train, y_train, X_test, y_test, epochs, learning_rate)

    # Evaluate the performance of the neural network
    predictions = np.argmax(nn.forward(X_test), axis=1)
```

```

accuracy = np.mean(predictions == y_test_labels)

return accuracy

# Create a study object and optimize the objective function
study = optuna.create_study(study_name='nn_study', direction='maximize')
study.optimize(objective, n_trials=10)

# Print the best hyperparameters
print(f"Best trial: {study.best_trial.params}")
print(f"Best value: {study.best_trial.value}")

```

```

[I 2024-03-25 21:25:03,832] A new study created in memory with name: nn_study
[I 2024-03-25 21:27:52,993] Trial 0 finished with value: 0.9666666666666667 and
parameters: {'hidden_size': 55, 'learning_rate': 0.043181890117968594, 'epochs':
7767}. Best is trial 0 with value: 0.9666666666666667.
[I 2024-03-25 21:30:29,852] Trial 1 finished with value: 0.21388888888888888 and
parameters: {'hidden_size': 53, 'learning_rate': 0.00024757383393530526,
'epochs': 9134}. Best is trial 0 with value: 0.9666666666666667.
[I 2024-03-25 21:32:09,488] Trial 2 finished with value: 0.16944444444444445 and
parameters: {'hidden_size': 97, 'learning_rate': 0.00038955617273236985,
'epochs': 4964}. Best is trial 0 with value: 0.9666666666666667.
[I 2024-03-25 21:32:26,650] Trial 3 finished with value: 0.13055555555555556 and
parameters: {'hidden_size': 91, 'learning_rate': 0.00016455567772291505,
'epochs': 940}. Best is trial 0 with value: 0.9666666666666667.
[I 2024-03-25 21:34:14,600] Trial 4 finished with value: 0.9444444444444444 and
parameters: {'hidden_size': 48, 'learning_rate': 0.028031573321476856, 'epochs':
7780}. Best is trial 0 with value: 0.9666666666666667.
[I 2024-03-25 21:36:26,633] Trial 5 finished with value: 0.3416666666666667 and
parameters: {'hidden_size': 124, 'learning_rate': 0.0004179701839858188,
'epochs': 7866}. Best is trial 0 with value: 0.9666666666666667.
[I 2024-03-25 21:39:21,789] Trial 6 finished with value: 0.825 and parameters:
{'hidden_size': 107, 'learning_rate': 0.0021982110768205094, 'epochs': 9456}.
Best is trial 0 with value: 0.9666666666666667.
[I 2024-03-25 21:40:33,457] Trial 7 finished with value: 0.16388888888888889 and
parameters: {'hidden_size': 100, 'learning_rate': 0.0005144087737501437,
'epochs': 3546}. Best is trial 0 with value: 0.9666666666666667.
[I 2024-03-25 21:41:00,385] Trial 8 finished with value: 0.14722222222222223 and
parameters: {'hidden_size': 94, 'learning_rate': 0.0016095589947516431,
'epochs': 1380}. Best is trial 0 with value: 0.9666666666666667.
[I 2024-03-25 21:42:52,334] Trial 9 finished with value: 0.875 and parameters:
{'hidden_size': 81, 'learning_rate': 0.006590500659615139, 'epochs': 6108}. Best
is trial 0 with value: 0.9666666666666667.

Best trial: {'hidden_size': 55, 'learning_rate': 0.043181890117968594, 'epochs':
7767}
Best value: 0.9666666666666667

```

1.10 Predict

```
[10]: best_nn = NeuralNetwork(input_size, study.best_trial.params['hidden_size'],  
    ↪output_size, loss_func)  
best_trainer = Trainer(best_nn, loss_func)  
best_trainer.train(X_train, y_train, X_test, y_test, study.best_trial.  
    ↪params['epochs'], study.best_trial.params['learning_rate'])  
  
# Evaluate the performance of the best neural network  
predictions = np.argmax(best_nn.forward(X_test), axis=1)  
accuracy = np.mean(predictions == y_test_labels)  
print(f"Best accuracy: {accuracy:.2%}")
```

Best accuracy: 96.67%